

```
In [1]: import pandas as pd
import numpy as np
from sklearn import metrics
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: from sklearn.datasets import load_boston
boston = load_boston()
```

```
In [4]: data = pd.DataFrame(boston.data)
```

```
In [5]: data.head()
```

Out[5]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [6]: #Adding the feature names to the dataframe
data.columns = boston.feature_names
data.head()
```

Out[6]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90

```
In [7]: #Adding target variable to dataframe
data['PRICE'] = boston.target
```

```
In [8]: #Check the shape of dataframe  
data.shape
```

```
Out[8]: (506, 14)
```

```
In [9]: data.columns
```

```
Out[9]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
               'PTRATIO', 'B', 'LSTAT', 'PRICE'],  
              dtype='object')
```

```
In [10]: data.dtypes
```

```
Out[10]: CRIM      float64  
ZN        float64  
INDUS     float64  
CHAS      float64  
NOX       float64  
RM        float64  
AGE       float64  
DIS       float64  
RAD       float64  
TAX       float64  
PTRATIO   float64  
B         float64  
LSTAT     float64  
PRICE     float64  
dtype: object
```

```
In [11]: # Identifying the unique number of values in the dataset  
data.nunique()
```

```
Out[11]: CRIM      504  
ZN        26  
INDUS     76  
CHAS      2  
NOX       81  
RM        446  
AGE       356  
DIS       412  
RAD        9  
TAX       66  
PTRATIO   46  
B         357  
LSTAT     455  
PRICE     229  
dtype: int64
```

In [12]: *# Check for missing values*
`data.isnull().sum()`

Out [12]:

CRIM	0
ZN	0
INDUS	0
CHAS	0
NOX	0
RM	0
AGE	0
DIS	0
RAD	0
TAX	0
PTRATIO	0
B	0
LSTAT	0
PRICE	0
	dtype: int64

In [13]: *# See rows with missing values*
`data[data.isnull().any(axis=1)]`

Out [13]:

CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
------	----	-------	------	-----	----	-----	-----	-----	-----	---------	---	-------	-------

In [14]: *# Viewing the data statistics*
`data.describe()`

Out [14]:

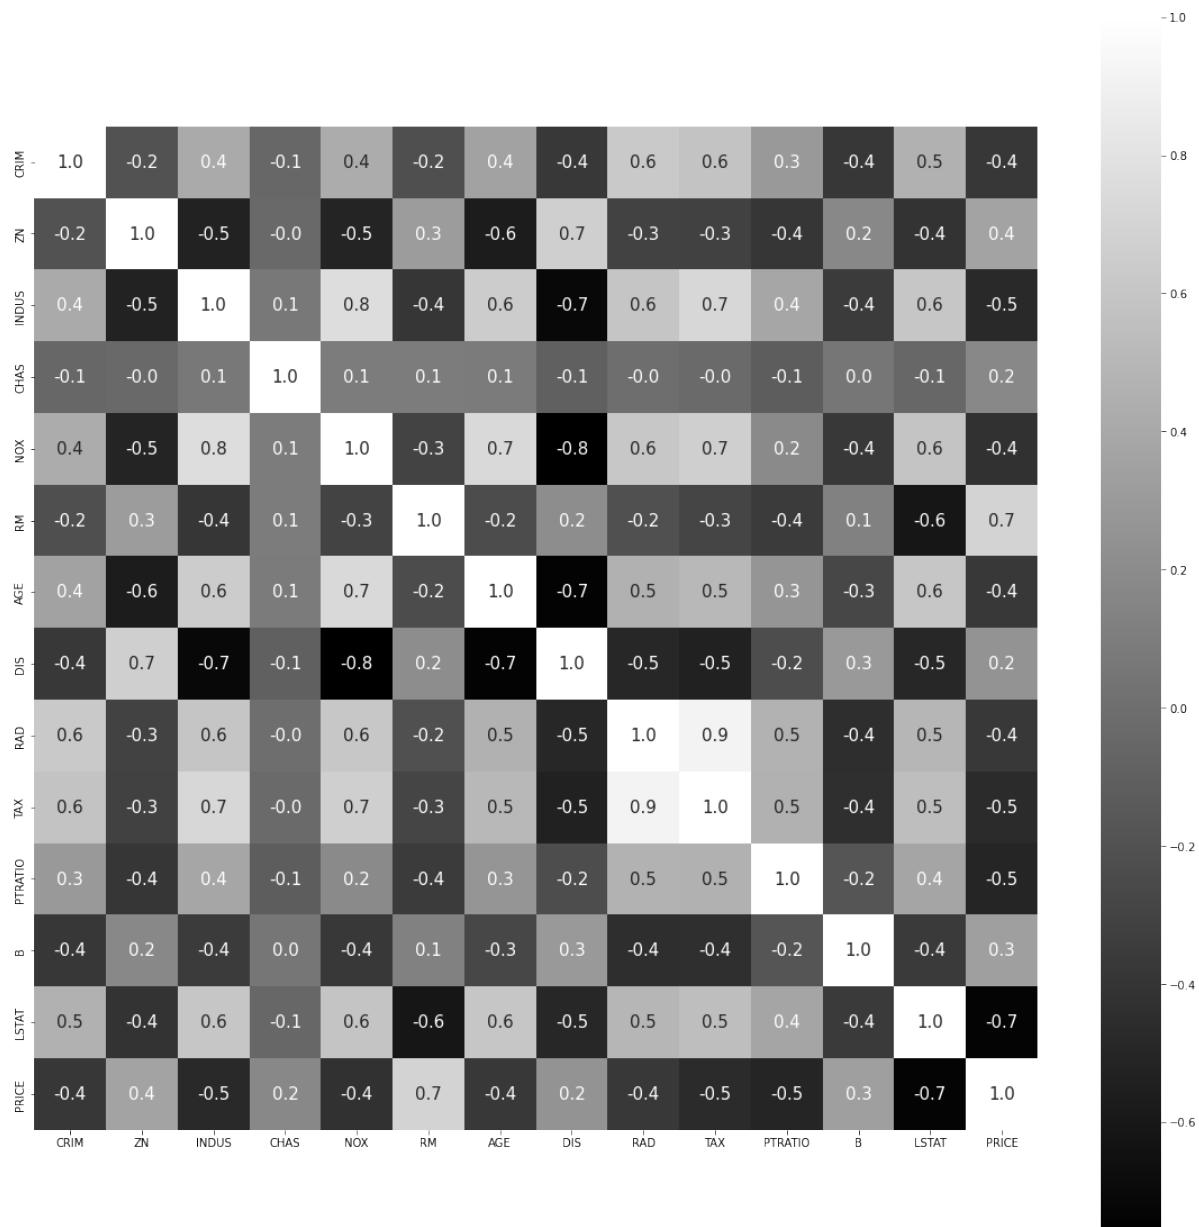
	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000

In [15]: *# Finding out the correlation between the features*
`corr = data.corr()`
`corr.shape`

Out [15]: (14, 14)

```
In [16]: # Plotting the heatmap of correlation between features
plt.figure(figsize=(20,20))
sns.heatmap(corr, cbar=True, square=True, fmt='.1f', annot=True, a
```

Out [16]: <AxesSubplot:>



```
In [17]: # Splitting target variable and independent variables
X = data.drop(['PRICE'], axis = 1)
y = data['PRICE']
```

In [18]: # Splitting to training and testing data

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size
```

```
In [19]: # Import library for Linear Regression
from sklearn.linear_model import LinearRegression
```

```
In [20]: # Create a Linear regressor
lm = LinearRegression()

# Train the model using the training sets
lm.fit(X_train, y_train)
```

```
Out[20]: LinearRegression()
```

```
In [21]: # Value of y intercept
lm.intercept_
```

```
Out[21]: 36.357041376594964
```

```
In [22]: #Converting the coefficient values to a dataframe
coefficients = pd.DataFrame([X_train.columns, lm.coef_]).T
coefficients = coefficients.rename(columns={0: 'Attribute', 1: 'Coeff'})
coefficients
```

```
Out[22]:
```

	Attribute	Coefficients
0	CRIM	-0.12257
1	ZN	0.055678
2	INDUS	-0.008834
3	CHAS	4.693448
4	NOX	-14.435783
5	RM	3.28008
6	AGE	-0.003448
7	DIS	-1.552144
8	RAD	0.32625
9	TAX	-0.014067
10	PTRATIO	-0.803275
11	B	0.009354
12	LSTAT	-0.523478

```
In [23]: # Model prediction on train data
y_pred = lm.predict(X_train)
```

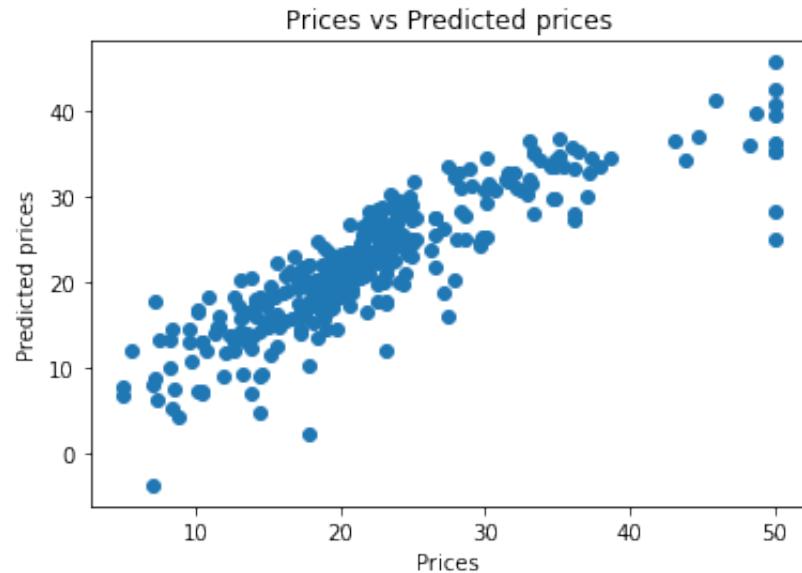
In [24]: *Evaluation*

```
R^2:',metrics.r2_score(y_train, y_pred))
Adjusted R^2:',1 - (1-metrics.r2_score(y_train, y_pred))*(len(y_train)-1)/(len(y_train)-1)
MAE:',metrics.mean_absolute_error(y_train, y_pred))
MSE:',metrics.mean_squared_error(y_train, y_pred))
RMSE:',np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

R^2: 0.7465991966746854
Adjusted R^2: 0.736910342429894
MAE: 3.0898610949711323
MSE: 19.07368870346903
RMSE: 4.367343437774162

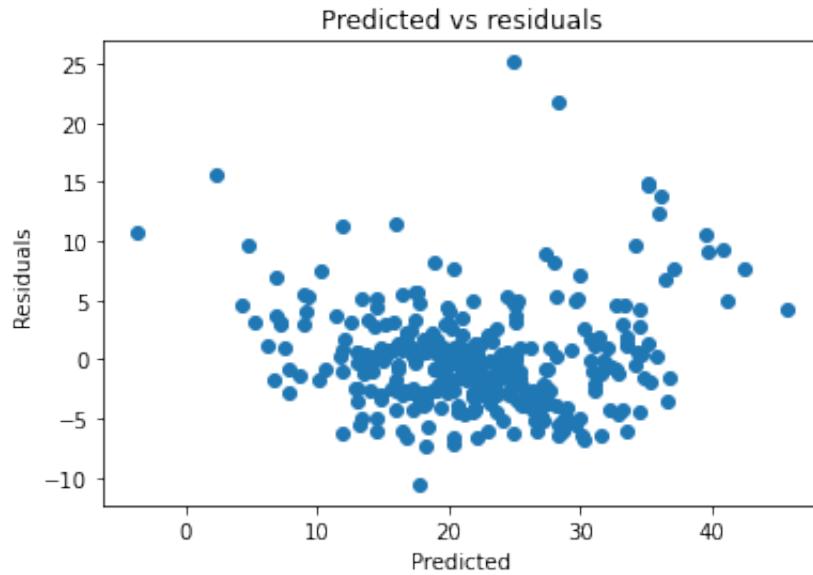
In [25]: *# Visualizing the differences between actual prices and predicted values*

```
plt.scatter(y_train, y_pred)
plt.xlabel("Prices")
plt.ylabel("Predicted prices")
plt.title("Prices vs Predicted prices")
plt.show()
```



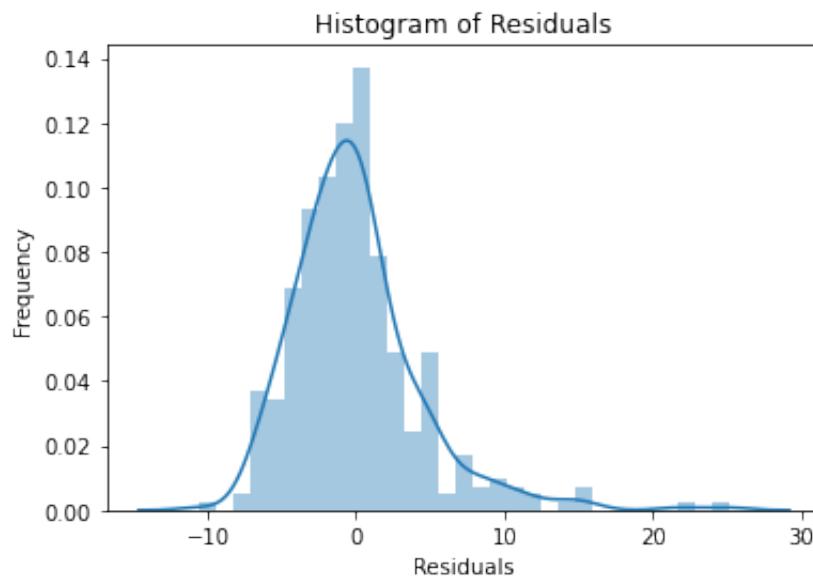
In [26]: *# Checking residuals*

```
plt.scatter(y_pred,y_train-y_pred)
plt.title("Predicted vs residuals")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.show()
```



In [27]: *# Checking Normality of errors*

```
sns.distplot(y_train-y_pred)
plt.title("Histogram of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```



In [28]:

```
# Predicting Test data with the model
y_test_pred = lm.predict(X_test)
```

In [29]:

```
linreg = metrics.r2_score(y_test, y_test_pred)
'R^2:', acc_linreg)
'Adjusted R^2:', 1 - (1-metrics.r2_score(y_test, y_test_pred))*(len(y
'MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
'MSE:', metrics.mean_squared_error(y_test, y_test_pred))
'RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

R^2: 0.7121818377409183
Adjusted R^2: 0.6850685326005699
MAE: 3.859005592370746
MSE: 30.053993307124266
RMSE: 5.482152251362987

In []: